

Blinded Random Corruption Attacks

Rodrigo Rubira Branco (@BSDaemon)
Principal Security Researcher
rodrigo.branco *noSPAM* intel.com

Shay Gueron (Core Architecture and Haifa University)
Senior Principal Engineer

DISCLAIMER

**We don't speak for our employer.
All the opinions and information
here are of our responsibilities**

**–So, mistakes and bad jokes are all
–OUR responsibilities**

Agenda

- Introduction
- The A-B-C Model for the Attacker's Level of Control and Targets
- Becoming “root” on a locked system with a BRBC attack
- Different attack scenarios and attack targets
- Mitigation Techniques
- Conclusions

Background

- Adversaries with physical access to attacked platform – are a concern
 - Mobile devices (stolen/lost)
 - Cloud computing (un-trusted environments)
- Read/write memory capabilities as an attack tool have been demonstrated:
 - Using different physical interfaces
 - Thunderbolt, Firewire, PCIe, PCMCIA and new USB standards
- Consequences of DRAM modification capabilities:
 - Active attacks on memory are possible
 - Attacker can change code / data **from any value to any chosen value**

Underlying attack assumption on the threat model:
The attacker has physical means to modify DRAM

Different attacker tactics

- Passive attack: the attacker can only eavesdrop DRAM contents, but is not able to inject or interfere with it (in-use or not)
 - Non-existent in reality
- Active static attack: the attacker can read DRAM contents but cannot modify in-use/to-be-used (saved) DRAM
 - Example: cold boot attack
 - The attack is on the data privacy
- Active dynamic attacks: the attacker can read and modify DRAM contents that are in-use/to-be-used (saved)

Memory Encryption without Authentication effectiveness limited to active static attacks since the ability to modify in-use/to-be-used DRAM is denied

Introduction (1)

- Some memory protection technologies against active dynamic attacks were proposed
 - Limiting the attacker's physical ability to read/write memory
 - E.g., blocking DMA access in some scenarios
 - Memory encryption
- **Memory encryption using “transparent encryption” mode:**
 - Simpler, cheaper, faster than “encryption + authentication”
 - Changes the assumptions on read/written memory capabilities of the attacker
 - Therefore, seems to be effective for limiting active dynamic attacks
- Memory encryption effects :
 - Attacker has **limited control** on the result of active attacks
 - But the physical memory modification **capabilities remain available**

Underlying attack assumption: attacker has physical means to modify DRAM

Introduction (2)

- Under memory encryption, the attacker has limited capabilities
 - **Blinded Random Block Corruption (BRBC)** attack
- **(Blinded)** The attacker does not know the plaintext memory values he can read from the (encrypted) memory.
- **(Random (Block) Corruption)** The attacker cannot control nor predict the plaintext value that would infiltrate the system when a modified (encrypted) DRAM value is read in and decrypted.
 - When using a block cipher (in standard mode of operation), any change in the ciphertext would **randomly corrupt** at least **one block** of the eventually decrypted plaintext
- The question: can memory encryption (that limits the active dynamic attacker capabilities to **BRBC** only) provide a “good enough” mitigation in practice?

Underlying attack assumption: attacker has physical means to modify DRAM

Introduction (3)

- We will show that:
 - Despite limited capabilities dynamic active attacks are still possible
 - Encryption-only does not offer a defense-in-depth mechanism against arbitrary memory overwrites **without removing capabilities assumptions**
- The BRBC attacker is able to create Time-of-check/Time-of-use (TOCTOU) race conditions all around the execution environment
 - Usual control-flow hijacking attacks require precise pointer control to redirect flow of execution. Usual DMA attacks perform precise code modification
 - Data-only attacks caused by a BRBC attacker can be induced after some code checks, therefore cause TOCTOU races that invalidate the results of such checks
 - Unexpected computation (and flows) can emerge (since code is driven by its input data)
 - Data-only based attacks, thus control flow enforcement can't prevent

Underlying attack assumption: attacker has physical means to modify DRAM

The A-B-C attacker model

- Access Seeking Attacker

This attacker is not the owner of the platform, but got it to his possession, in a locked state. He wishes to get an user access, in order to steal the data on the system.

- Breaching Attacker


This attacker is a legitimate user of the platform, who wishes to breach some of the system's policies or circumvent restrictions on his privileges.

- Conspirator Attacker


This attacker is also a legitimate user of the platform/environment. He has administrative powers and conspires to collect other users' data.

Underlying attack assumption: attacker has physical means to modify DRAM

Becoming “root” on a locked system with a BRBC attack

```
global var1...varn ←   
global preauth_flag  
global preauth_related...  
code_logic() {  
    if (preauth_enabled) {  
        call_preauth_mechanism() -> sets preauth_flag if successful  
    }  
repeat_auth:  
    if (preauth_flag) goto auth_ok;  
  
    authentication_logic();  
  
auth_ok:  
    return;  
}
```


Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag 
global preauth_related...
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();

auth_ok:
    return;
}
```


Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related... 
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();

auth_ok:
    return;
}
```

Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related...
code_logic() { 
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();


auth_ok:
    return;
}
```

Becoming “root” on a locked system with a BRBC attack


```
global var1...varn
global preauth_flag
global preauth_related...
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();

auth_ok:
    return;
}
```



Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related...
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth: 
    if (preauth_flag) goto auth_ok;

    authentication_logic();

auth_ok:
    return;
}
```

Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related...
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;
    authentication_logic();


    auth_ok:
        return;
}
```

← BRBC Attack to the preauth_flag

Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related...
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();           -> THIS NEVER GETS EXECUTED!

auth_ok: 
    return;
}
```

TOCTOU (Time-of-use/Time-of-check) Race Condition

- This was caused by our arbitrary memory write (the BRBC)
- The corrupted values adjacent to the `preauth_flag` were not used at this moment (thus the block corruption is not a problem)
- The check for the `preauth_flag` only checks for not 0 (thus we don't need to control the exact value)
- But how do we win the race?
 - In this case, quite simple: We just cause the authentication to fail at the first time (when it does ask the password)
 - The system waits for the password prompt
 - We cause the corruption and input invalid password
 - The authentication fails and the logic is repeated, but this time with the corruption!

Experiment

- The demonstration was created focusing on the underlying attack assumption
 - We use a debugger to make it easy to step through and see the corruption effect
 - We use the JTAG to demonstrate the physical addresses are not a concern
- SW mitigations are not feasible because the attacker has lots of possibilities for targets (not only ! 0 comparisons). Some examples:
 - If an attacker overwrites the NULL terminator of a string, he can generate buffer overflows, memory leaks
 - If an attacker overwrites an index, he can generate out-of-bounds writes, that might lead to user-mode dereferences if in kernel-mode context
 - If an attacker overwrites a counter, he can generate REFCOUNT overflows, leading to use-after-free conditions

Underlying attack assumption: attacker has physical means to modify DRAM

Attack Demonstration Using a Debugger

```
Debian GNU/Linux 7 devel tty2  
devel login: root  
Password: _
```

Attack Demonstration Using a Debugger

```
root(tty1)@devel:~/Shay# ps ax |grep login
2987 tty2      Ss+      0:00 /bin/login --
2989 tty1      S+       0:00 /bin/login
root(tty1)@devel:~/Shay# gdb /bin/login 2987
GNU gdb (GDB) 7.4.1-debian
```

Attack Demonstration Using a Debugger

```
0xb77cf424 in kernel_vsyscall ()  
(gdb) b *0x804a6e6  
Breakpoint 1 at 0x804a6e6: file login.c, line 966.  
(gdb) c  
Continuing.
```

Attack Demonstration Using a Debugger

```
Debian GNU/Linux 7 devel tty2
```

```
devel login: root
```

```
password:
```

```
login incorrect
```

```
devel login: root
```

Attack Demonstration Using a Debugger

```
0xb7711424 in __kernel_vsyscall ()
(gdb) b *0x804a6e6
Breakpoint 1 at 0x804a6e6: file login.c, line 966.
(gdb) c
Continuing.

Breakpoint 1, 0x0804a6e6 in main (argc=3, argv=0xbfe82554) at login.c:966
966                                spwd = xgets_pnam (username);
(gdb) set preauth_flag="16 bytes garbage"
(gdb) c
Continuing.
process 2987 is executing new program: /bin/bash
```


Attack Demonstration Using a Debugger

```
Debian GNU/Linux 7 devel tty2


devel login: root
Password:
Login incorrect

devel login: root
1 failure since last login.
Last was Mon Mar  9 11:31:26 2015 on /dev/tty2.
root(tty2)@devel:~# whoami
root
root(tty2)@devel:~# id
uid=0(root) gid=0(root) groups=0(root)
root(tty2)@devel:~# _
```

Attack Demonstration using the JTAG Interface

- The difference on the JTAG demonstration is:
 - Establish the possibility of the attack against the physical address space instead of the virtual one (as with the debugger)
 - Demonstrate that blinded reads are enough to gather locality of the targeted overwrite
 - Understand possible mitigations and their impacts on the attack (for example, control-flow enforcement technologies would not have prevented the attack either and can't be considered another layer of defense against BRBC)
- Limitations of the JTAG attack
 - For the MEE case, the JTAG access would be encrypted/decrypted, thus it would not be dealing with the encrypted content

Attack Demonstration using the JTAG Interface



Different Attack Scenarios and Attack Targets

- Attacker with user privileges on the machine
 - Higher control/visibility of the memory space
 - Tries to bypass security policies
 - Local administrator (common on cloud-based scenarios)
- All system software/components can be seen as targets
 - We just demonstrated in a highly-limited scenario (locked machine, unknown software running, little to no information on the OS details)
- As more interactions with the system, as bigger is the scope of possible attack targets (as discussed previously)

Mitigation Techniques

- Hibernation when used together with proper disk encryption
- VT-d/IOMMU and PMRs
 - Limits DMA capabilities exposed
 - Might not be enough against certain attackers (that have physical access) and in some platforms (only effective if the attack requirement is fully removed)
- Software self-protection (or control flow enforcement technologies)
 - Attack uses valid flows with invalid data (data-only attack) bypassing them
 - Different attack targets make software hardening inviable
- Memory encryption with Authentication
 - Able to detect the arbitrary change and prevent the attack
- Intel SGX (Software Guard eXtensions)
 - Currently employ authentication and replay protection

Conclusions

- Formalization of the BRBC attack
- Hierarchical model of the A-B-C attackers
- Practical demonstration of a BRBC attack
 - Definition of premises
 - Data-only attack generating a TOCTOU
- Discussion on mitigation paths
- **Encryption-only by itself is not necessarily a “good enough” defense-in-depth mechanism against arbitrary memory write primitive**

To consider

- Since encryption-only is not a defense against arbitrary writes (that not only break integrity, but also the confidentiality as demonstrated):
 - What is easier/viable: Remove ***ALL*** cases of arbitrary writes for ***ALL*** platforms the technology would support (which would depend on integration teams capabilities to guarantee that) or
 - At the technology level support encryption with authentication, which is a solution based only on itself

End! Really is !?

Rodrigo Rubira Branco (@BSDaemon)

Principal Security Researcher – Security Center of Excellence Core Client
rodrigo.branco *noSPAM* intel.com

Shay Gueron

Senior Principal Engineer - Core Architecture and Haifa University